

Persisted-cache cold start to attend-ready (warm page-cache read + H2D + reconstruct), Llama-3.1-8B, A100-40GB, medians of 5

